
Chapter 6

Congestion Control and Resource Allocation


Congestion Control and Resource Allocation

- The issue is “how to **effectively** and **fairly** allocate resources among a collection of competing users?”
- The resources being shared include
 - The **bandwidth** of the links
 - The **buffers** on the routers or switches
- **Congested:** when too many packets are contending for the same link, the **queue overflows** and packets have to be **dropped frequently**
 - A **congestion-control mechanism** is introduced
- Congestion may be avoided by using good **resource allocation mechanism**
 - Possibly make congestion-control unnecessary

Issues in Resource Allocation


Issues in Resource Allocation

- The difference between **flow control** and **congestion control**
 - Flow control involves keeping **a fast sender** from overrunning a slow **receiver**
 - Congestion control is intended to keep **a set of senders** from sending too much data into **the network**




Network Model

- A packet-switched network consists of multiple links and routers (or switches)
- A given source may have **more than enough capacity** on the **immediate** outgoing link to send a packet
- In the middle of a network, the packets encounter a link that is being used by many **different traffic sources**
 - The congested router is sometimes call the **bottleneck**



Connectionless Flows

- For the Internet, the network is essentially **connectionless**, with any **connection-oriented** service implemented in the transport protocol that is running on the end hosts
 - IP** provides a **connectionless** datagram delivery service
 - TCP** implements an end-to-end **connection abstraction**
- Flow**: a sequence of packets sent between a **source/destination** pair and following **the same route** through the network



Connectionless Flows


- Each router maintains **soft state and hard state information** for each flow
- **Soft state** is **not** explicitly created and removed by **signaling**
 - That can be used to make **resource allocation** decisions about the packets that belong to the flow
- **Hard state** is explicitly created and removed by **signaling**
 - That is generally used to make the packets being **correctly routed** from the source to the destination
- The correct operation of the network does not depend on soft state information, but the router with this information can **better handle the packets**
 - Resource allocation and congestion control

Router-Centric Versus Host-Centric

- The **Router-Centric** mechanism addresses the problem from **inside** the network
 - Each **router** takes responsibility for deciding **when** packets are forwarded and selecting **which** packets are dropped
- The **Host-Centric** mechanism addresses the problem from the **edges** of the network
 - The **end hosts** observe the network conditions and adjust their behavior accordingly
- These two groups are **not mutually exclusive**
 - A Router-Centric network still expects the end hosts to adhere to any **advisory messages sent by the routers**
 - A Host-Centric network still has some policy for deciding **which packets to drop** when their queues do overflow


Reservation-Based

- In a **Reservation-Based** system, the end host asks for a certain amount of capacity at the time a flow is established
 - Each router allocates **enough resources** to satisfy the request
 - If the request cannot be satisfied at some router, the router **rejects** the flow



Feedback-Based

- In a **Feedback-Based** system, the end hosts begin sending data **without reserving any capacity** and **adjust their sending rate** according to the received feedback
 - **Explicit feedback**: it is according to a message from a congested router
 - **Implicit feedback**: it is according to the **externally observable behavior** of the network (e.g. **packet losses**)



Reservation-Based Versus Feedback-Based

- A **Reservation-Based** system always implies **a router-centric** resource allocation mechanism
 - Each router is responsible for keeping track the resource
- A **Feedback-Based** system can imply **either a router- or host-centric** mechanism
 - If the feedback is **explicit**, the router is involved
 - If the feedback is **implicit**, the routers **silently drop packets** when they become congested

Window-Based Versus Rate-Based


- Both **flow-control** and **resource allocation** mechanisms need a way to express to the sender
- For **Window-Based** mechanism (such as TCP), the receiver advertises a window to the sender
 - This window corresponds to how much **buffer space** the receiver has
- For **Rate-Based** mechanism, a sender's behavior is controlled by using a **rate**
- **Rate-based** characterization of flows is a good choice in a reservation-based system
 - Supports different qualities of service (**QoS**)

Evaluation Criteria (Effectiveness)

- A network should **effectively** and **fairly** allocate its resources
- For evaluating effectiveness, two metrics of networking are considered: **throughput** and **delay**
- One way to increase throughput is to allow **as many** packets into the network **as possible**
 - Drive the utilization of all the links **up to 100%**
 - Increase the length of the **queues** at each router
 - Longer queues mean packets are **delayed longer** in the network \Rightarrow a large delay
- We may use the ratio of throughput to delay as a metric for evaluating the effectiveness of a resource allocation scheme
 - **Power = Throughput / Delay**

Evaluation Criteria (Effectiveness)

- Power is based on an **M/M/1 queuing network** that assumes **infinite queues**
 - Real networks have only **finite buffers**
- Power is typically defined relative to a **single connection**
 - Real networks have **multiple, competing connections**
- The objective is to **maximize** the Power ratio
 - Ideally, the resource allocation mechanism would operate at the **peak** of this curve




Evaluation Criteria (Fairness)

- A **reservation-based** resource allocation scheme provides an explicit way to create **controlled unfairness**
- In the **absence** of explicit information, we would like for each flow to receive an **equal share** of the bandwidth
- But **equal shares** may not equate to **fair shares**
 - Should we consider the length of the paths being compared?

If the same bandwidth is allocated,

Is it fair?



Evaluation Criteria (Fairness)

- **Raj Jain** has proposed a metric that can be used to **quantify the fairness** of a congestion control mechanism
- Given a set of flow throughputs **(x_1, x_2, \dots, x_n)** bits/second
- A **fairness index function** is defined as

$$f(x_1, x_2, \dots, x_n) = \frac{(\sum_{i=1}^n x_i)^2}{n \sum_{i=1}^n x_i^2}$$

- Between 0 and 1, with **1 representing greatest fairness**
- All **n** flows receive a throughput of 1 unit of data per second
 - **$x_1 = x_2 = \dots = x_n$**

$$f(x_1, x_2, \dots, x_n) = n^2 / n \times n = 1$$

Evaluation Criteria (Fairness)

- **Fairness:** all flows have the same throughput

- $x_1 = x_2 = \dots = x_n$

- Suppose one flow receives a throughput of $1+\Delta$



$$f(x_1, x_2, \dots, x_n) = (n^2 + 2n\Delta + \Delta^2) / (n^2 + 2n\Delta + \Delta^2) < 1$$

- No matter Δ is positive or negative, the index **drops** below 1

Queuing Disciplines

FIFO (First-In-First-Out)

- The idea of **FIFO** queuing is “The first packet that arrives at a router is the first packet to be transmitted”
 - Also called **first-come-first-served (FCFS)** queuing

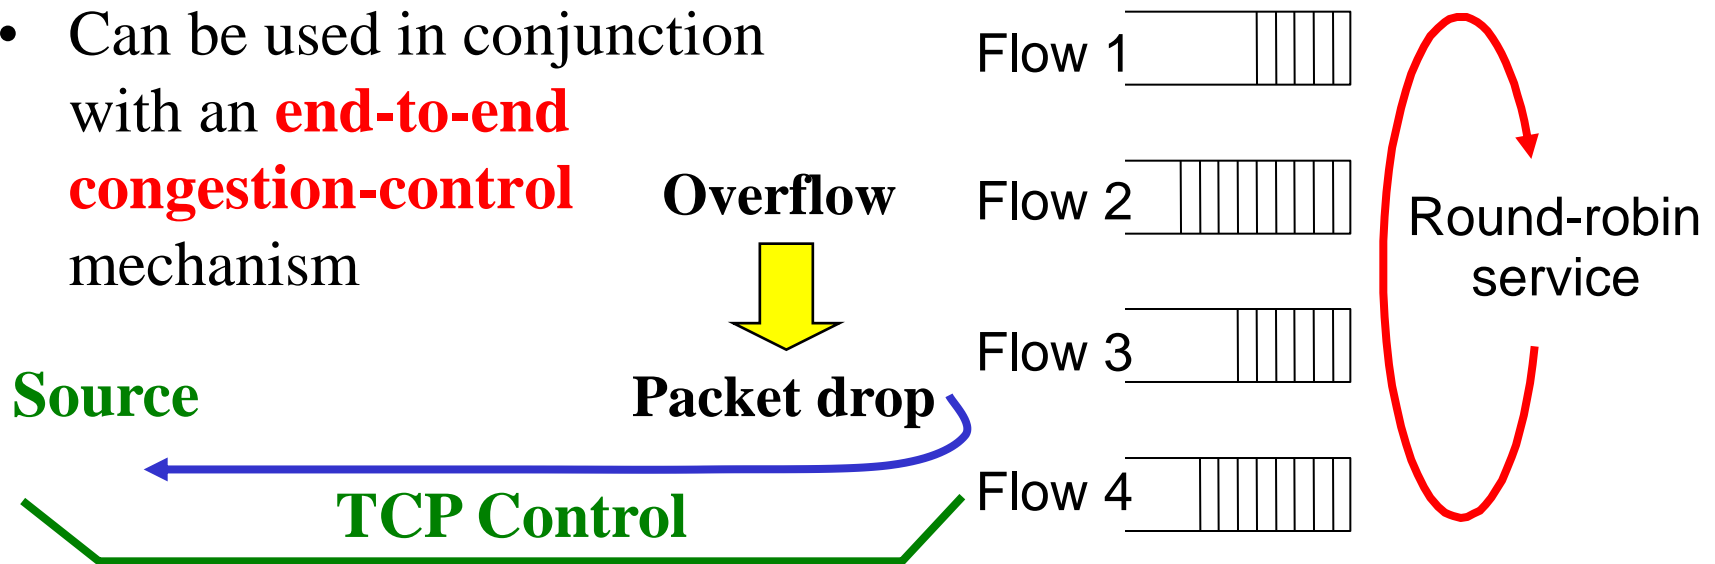


FIFO

- FIFO queuing is “**FIFO with tail drop**” (most widely used)
 - **FIFO** (scheduling discipline): determines the **order** in which packets are transmitted
 - **Tail drop** (drop policy): determines which packets get dropped
- A simple variation on basic FIFO queuing is **priority queuing**
 - Make each packet with priority (carried in the IP **Type of Service (TOS)** field)
 - The router always transmits packets out of the **highest-priority queue** if the queue is nonempty
 - Then moves on to the **next priority queue**
 - Within each priority, packets are still **FIFO**

FQ (Fair Queuing)

- The main problem with FIFO queuing is that it **does not discriminate** between **different traffic sources**
- **FQ** maintains a **separate queue** for each flow currently being handled by the router
 - Services these queues in a **round-robin** manner
- Can be used in conjunction with an **end-to-end congestion-control** mechanism




FQ (Fair Queuing)

- The packets being processed at a router are not necessarily the same length
 - Takes **packet length** into consideration
- Let P_i denote the length of packet i , S_i denote the time when the router starts to transmit packet i , and F_i denote the time when the router finishes transmitting packet i
 - **$F_i = S_i + P_i$**
- Let A_i denote the time that packet i arrives at the router
 - **$S_i = \max(F_{i-1}, A_i)$**
 - **$F_i = \max(F_{i-1}, A_i) + P_i$**

FQ (Fair Queuing)

- All the F_i are treated as **timestamps**
 - The next packet to transmit is the packet with **lowest** timestamp
 - However, a **newly** arriving packet **cannot preempt** a packet that is currently being transmitted
- If the link is fully loaded and there are n flows, each flow shares **$1/n$ of the link bandwidth** (Not perfect: can't preempt current packet)



Shorter packets are sent first

Sending of longer packet is completed first

TCP Congestion Control

TCP Congestion Control

- **TCP** applies an **end-to-end** congestion control mechanism
- The essential strategy of TCP is to send packets into the network **without a reservation**
 - To react to **observable events** that occur
- The idea of TCP congestion control is for each source to determine how much **capacity** is available in the network
- By using **ACKs** to pace the transmission of packets
 - TCP is said to be **self-clocking**

Additive Increase/Multiplicative Decrease


- TCP maintains a new state variable for each connection
 - Called **CongestionWindow**
 - Used by the **source** to limit how much data it is allowed to have in transit at a given time
- The TCP's **effective window**:
 - **MaxWindow =**
MIN (CongestionWindow, AdvertisedWindow)
 - **EffectiveWindow =**
MaxWindow – (LastByteSent–LastByteAcked)

Flow Control (Advertised Window)

- **AdvertiseWindow** is sent by the receiver
- To avoid overflowing the receive buffer, the sender computes an effective window that limits how much data it can send:

– **EffectiveWindow** = (flow control)

$$\text{AdvertiseWindow} - (\text{LastByteSent} - \text{LastByteAked})$$



– **EffectiveWindow** = (congestion control)

$$\text{MaxWindow} - (\text{LastByteSent} - \text{LastByteAked})$$

- **CongestionWindow \geq AdvertiseWindow**: original “Sliding Window Algorithm”
- **CongestionWindow $<$ AdvertiseWindow**: control by network congestion

Additive Increase/Multiplicative Decrease

- How TCP learns an appropriate value for **CongestionWindow**
 - The TCP source sets the **CongestionWindow** based on the level of congestion in the network
 - **Decrease** the congestion window when congestion goes up
 - **Increase** the congestion window when congestion goes down
- The mechanism is called **Additive Increase/Multiplicative Decrease (AIMD)**

Additive Increase/Multiplicative Decrease

- How does the source determine that the network is **congested**?
 - The main reason of packets are not delivered, i.e. **timeout**, is that packet were dropped **due to congestion**
- TCP interprets timeouts as **a sign of congestion** and then **reduces** the transmission rate
- **Multiplicative Decrease:**
 - Each time a timeout occurs, the source sets **CongestionWindow** to **half** of its previous value
 - **CW** = 8; timeout \Rightarrow **CW** = 4; timeout \Rightarrow **CW** = 2; timeout \Rightarrow **CW** = 1
 - **CongestionWindow** is **not allowed** to fall below the size of a single packet, i.e. the **maximum segment size (MSS)**

Additive Increase/Multiplicative Decrease

- It also needs to increase the **CongestionWindow** to take advantage of **newly available** bandwidth in the network

- Additive Increase:**


- Every time the source **successfully** sends a CongestionWindow's worth of packets, it **adds** the equivalent of one packet to CongestionWindow
- Specifically, each time an ACK arrives the congestion window is incremented

$\text{Increment} = \text{MSS} \times (\text{MSS} / \text{CongestionWindow})$

$\text{CongestionWindow} += \text{Increment}$


$\text{Fraction} = \text{MSS} / \text{CongestionWindow}$

Source Destination




Additive Increase/Multiplicative Decrease

- This pattern of continually increasing and decreasing the congestion window continuous throughout the connection
 - **A sawtooth pattern**
- The source is willing to **reduce** its congestion window at a **much faster** rate than it is willing to **increase** this window
 - The consequences of having too large a window are much worse than those of it being too small




AIMD Fairness



Slow Start

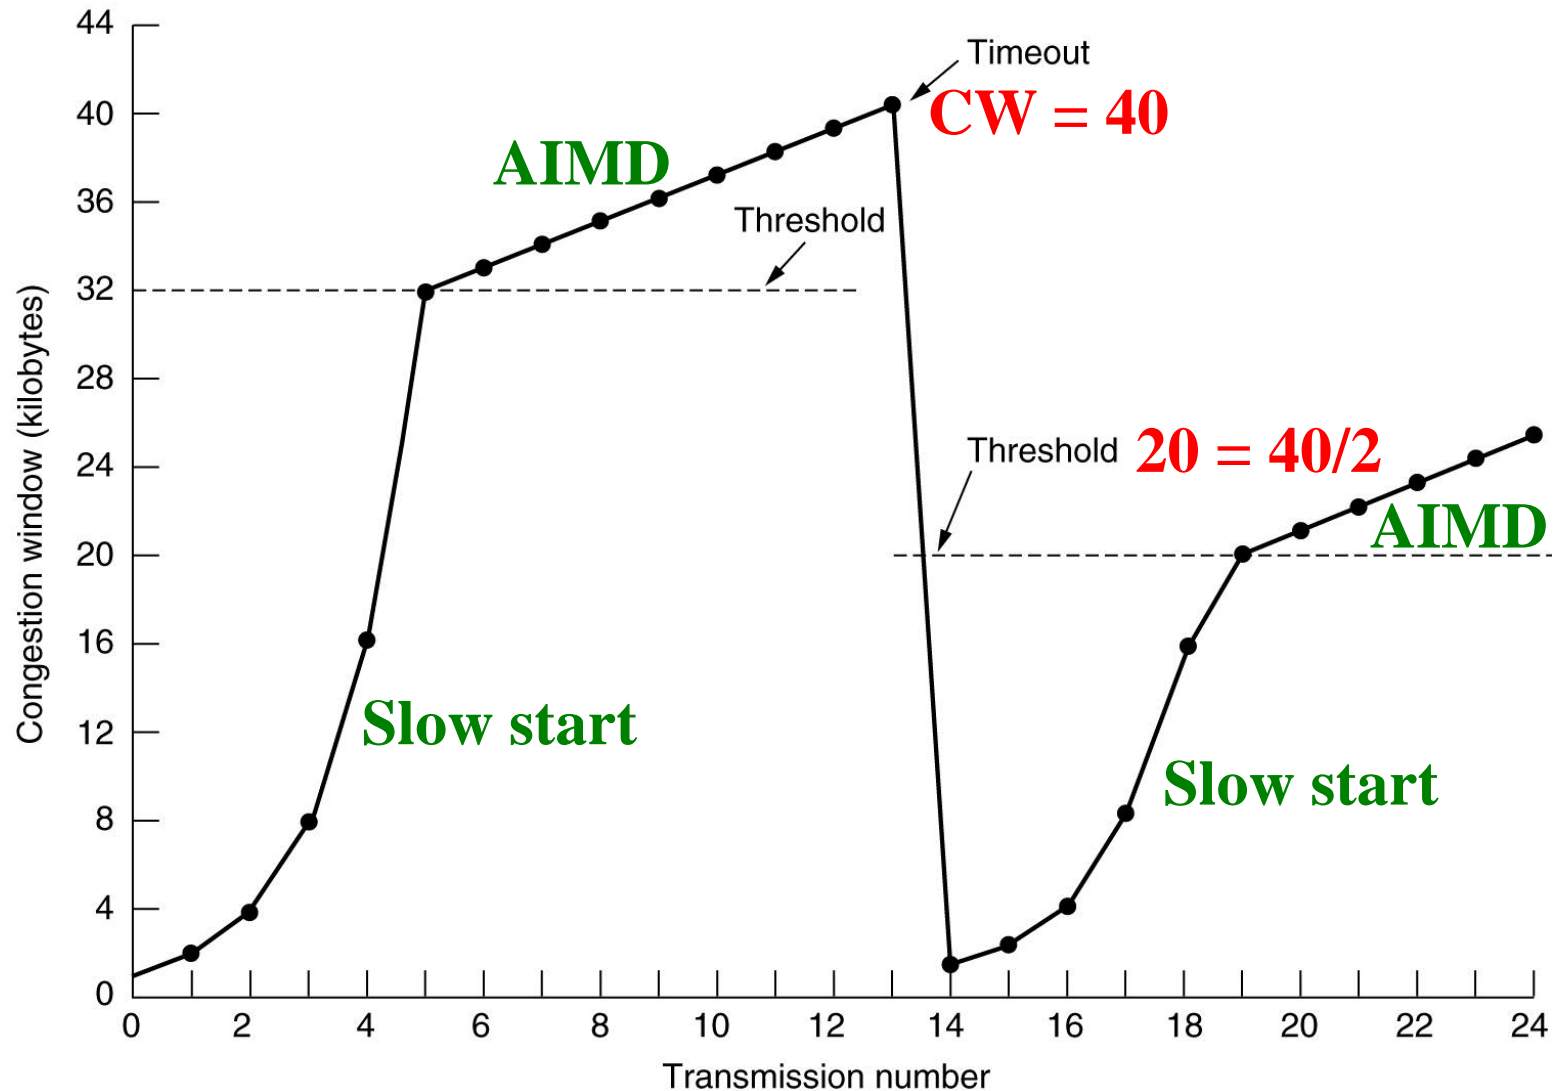
- The AIMD **takes too long to ramp up** a connection from its start to the available bandwidth
 - TCP provides a **second mechanism**, called **slow start**, that is used to increase the congestion window **rapidly** from a cold start
 - **Exponentially** rather than linearly
- For each received ACK, TCP increments Congestion Window by **2**
 - More **rapidly** than **Additive Increase**



Slow Start


- Two different situations in which slow start runs
- The **very beginning** of a connection
 - The source has **no idea** about the available bandwidth
 - **Double** congestion window for each RTT until **there is a packet loss**
- The connection **goes dead** while **waiting for a timeout to occur** (no ACK comes back, and no packet can be transmitted)
 - The source uses **slow start** to restart the flow of data
 - **Target congestion window** (**CongestionThreshold, CT**)
 - Set to the value of CongestionWindow, that existed prior to the last packet loss, divided by 2
 - After CongestionWindow **has reached the target**, the additive increase (AIMD) is used beyond this point

Slow Start



Slow Start (cont)

- Exponential growth, but slower than all at once
- Used...
 - when first starting connection
 - when connection goes dead waiting for timeout
- Trace




- Problem: lose up to half a **CongestionWindow**'s worth of data

Fast Retransmit

- TCP timeouts lead to long periods of time during which the connection went dead while **waiting for a timer to expire**
 - The **fast retransmit mechanism** was added to TCP
 - Triggers the retransmission of a **dropped packet** sooner than the regular timeout mechanism
- Every time a data packet arrives at the receiving side, the receiver responds with a **acknowledgment**
 - If a packet arrives out of order, TCP resends the **last ACK**
 - This **duplicate ACK** suggests that an earlier packet might have been **lost** (it is possible that it just only be delayed)
- In practice, TCP waits until it has seen **three** duplicate ACKs before retransmitting the packet

Fast Retransmit

- The destination receives packets 1 and 2
- Packet 3 is lost in the network
- When packet 4 is received, the destination resends ACK2
- ...
- When three duplicate ACK2 is received by the sender
 - **Retransmit packet 3**
- When packet 3 is received, the destination sends a cumulative ACK up to packet 6




Fast Retransmit


- For TCP with **fast retransmit** mechanism
 - The long period which the congestion window **stays flat** and no packets are sent **have been eliminated**
 - Result in roughly a **20% improvement** in the throughput

Fast Recovery

- **Fast recovery:** uses the ACKs that are still in the pipe to clock the sending of packets
 - **Remove the slow start phase** between when fast retransmit detects a lost packet and additive increase begins
 - Simply **cuts the congestion window in half** and **resumes additive increase**



TCP Equation



- Total number of packets transmitted in a period

$$= \frac{w}{2} + \left(\frac{w}{2} + 1\right) + \left(\frac{w}{2} + 2\right) + \dots + w \approx \frac{3}{8} w^2$$

- One packet dropped in a period

$$\rho \approx \frac{1}{\frac{3}{8} w^2} \Rightarrow w \approx \sqrt{\frac{81}{3\rho}}$$


$$\bullet \text{RATE} \approx \frac{\frac{3}{8} w^2}{\frac{w}{2} RTT} \approx \frac{3}{4} \frac{w}{RTT} \approx \frac{1}{RTT} \frac{1}{\sqrt{\rho}} \sqrt{\frac{3}{2}} \approx \frac{1.22}{RTT \sqrt{\rho}}$$

Congestion Avoidance

- TCP's strategy
 - control congestion once it happens
 - repeatedly increase load in an effort to find the point at which congestion occurs, and then back off
- Alternative strategy
 - predict when congestion is about to happen
 - reduce rate before packets start being discarded
 - call this congestion *avoidance*, instead of congestion *control*
- Two possibilities
 - router-centric: DECbit and RED Gateways
 - host-centric: TCP Vegas

DECbit

- Add binary congestion bit to each packet header
- Router
 - monitors average queue length over last busy+idle cycle



- set congestion bit if average queue length > 1
- attempts to balance throughput against delay

End Hosts

- Destination echoes bit back to source
- Source records how many packets resulted in set bit
- If less than 50% of last window's worth had bit set
 - increase **CongestionWindow** by 1 packet
- If 50% or more of last window's worth had bit set
 - decrease **CongestionWindow** by 0.875 times

Random Early Detection (RED)

- Notification is implicit
 - just drop the packet (TCP will timeout)
 - could make explicit by marking the packet
- Early random drop
 - rather than wait for queue to become full, drop each arriving packet with some *drop probability* whenever the queue length exceeds some *drop level*


RED Details

- Compute average queue length

$$\text{AvgLen} = (1 - \text{Weight}) * \text{AvgLen} + \text{Weight} * \text{SampleLen}$$

$0 < \text{Weight} < 1$ (usually 0.002)

SampleLen is queue length each time a packet arrives



RED Details (cont)

- Two queue length thresholds

```
if AvgLen <= MinThreshold then
```

```
    enqueue the packet
```

```
if MinThreshold < AvgLen < MaxThreshold then
```

```
    calculate probability P
```

```
    drop arriving packet with probability P
```

```
if MaxThreshold <= AvgLen then
```

```
    drop arriving packet
```


RED Details (cont)

- Computing probability P

$$\text{TempP} = \text{MaxP} * (\text{AvgLen} - \text{MinThreshold}) / (\text{MaxThreshold} - \text{MinThreshold})$$

$$P = \text{TempP} / (1 - \text{count} * \text{TempP})$$

- Drop Probability Curve



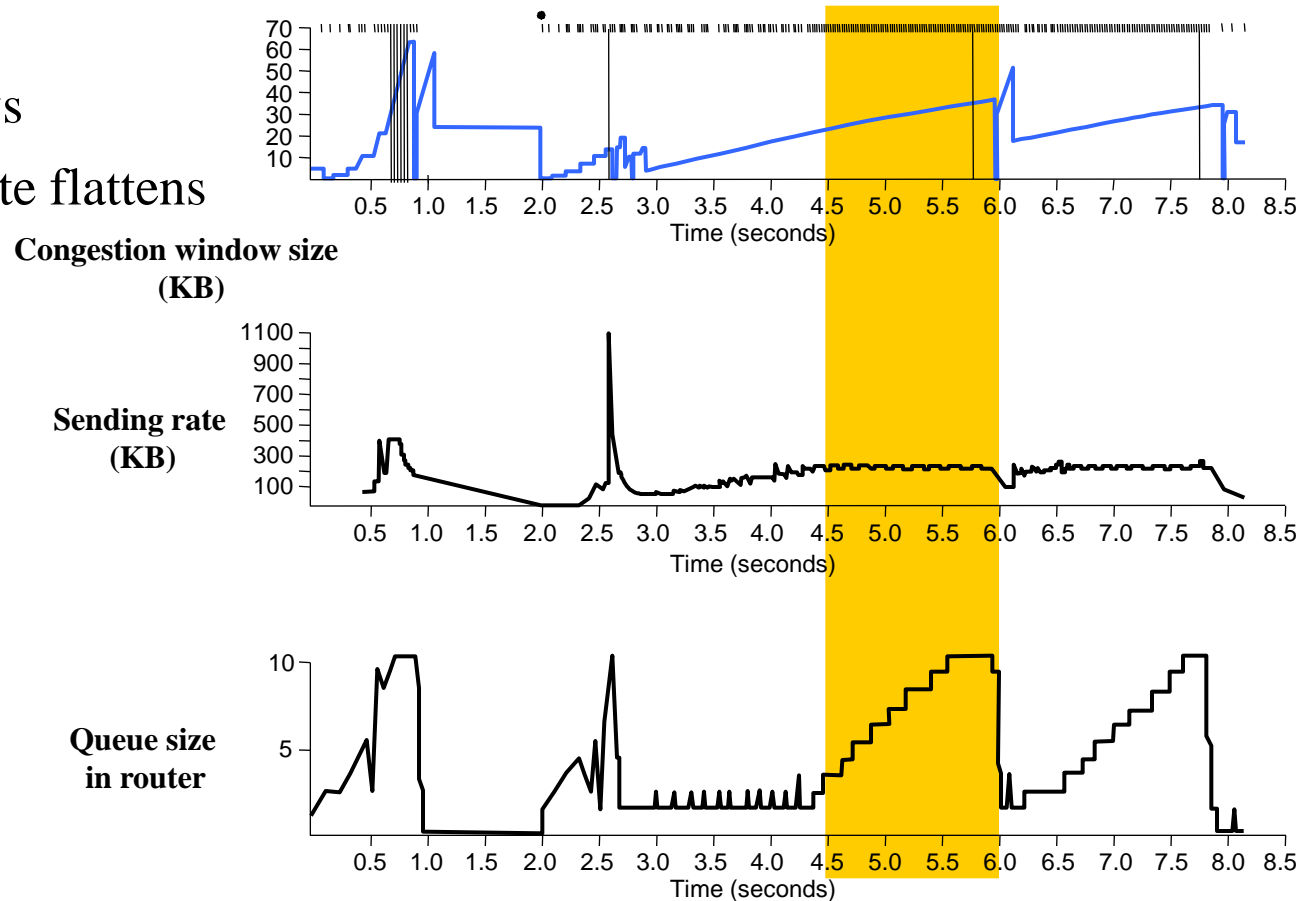
Tuning RED

- Probability of dropping a particular flow's packet(s) is roughly proportional to the share of the bandwidth that flow is currently getting
- **MaxP** is typically set to 0.02 , meaning that when the average queue size is halfway between the two thresholds, the gateway drops roughly one out of 50 packets.
- If traffic is bursty, then **MinThreshold** should be sufficiently large to allow link utilization to be maintained at an acceptably high level
- Difference between two thresholds should be larger than the typical increase in the calculated average queue length in one RTT; setting **MaxThreshold** to twice **MinThreshold** is reasonable for traffic

TCP Vegas

- Idea: source watches for some sign that router's queue is building up and congestion will happen too; e.g.,

- RTT grows
- sending rate flattens



Algorithm

- Let **BaseRTT** be the minimum of all measured RTTs (commonly the RTT of the first packet)
- If not overflowing the connection, then
$$\text{ExpectRate} = \text{CongestionWindow} / \text{BaseRTT}$$
- Source calculates sending rate (**ActualRate**) once per RTT
- Source compares **ActualRate** with **ExpectRate**

Diff = **ExpectedRate** - **ActualRate**

if **Diff** < α

 increase **CongestionWindow** linearly

else if **Diff** > β

 decrease **CongestionWindow** linearly


else

 leave **CongestionWindow** unchanged



Algorithm (cont)

- Parameters
 - $\alpha = 1$ packet
 - $\beta = 3$ packets

Congestion window size



- Expected throughput
- Actual throughput
- Region between α & β



Quality of Service

Quality of Service

- Packet-switched networks have offered the promise of supporting **multimedia applications**
 - One obstacle is the need for higher-bandwidth links
 - Improvements in coding have reduced the bandwidth needs
- **Audio or video services** are **delay sensitive** (real-time applications). \Rightarrow **timely delivery** is very important
 - Need some sort of assurance from the network that data is likely to **arrive on time**
- **Non-real-time applications:** use an end-to-end **retransmission** strategy to make sure that data arrives **correctly**
 - Such a strategy cannot provide timeliness
 - Retransmission \Rightarrow **long latency**

Quality of Service

- **Timely arrival** must be provided by the **network** itself (routers), not just at the network edges (hosts)
- **Best-effort model** is **not sufficient** for real-time applications
 - A new service model is required
 - Applications that need **higher assurances** can ask the network for them
- The network will treat some packets (**real-time**) differently from others (**non-real-time**)
- A network that can provide these different levels of service is often said to support **quality of service (QoS)**

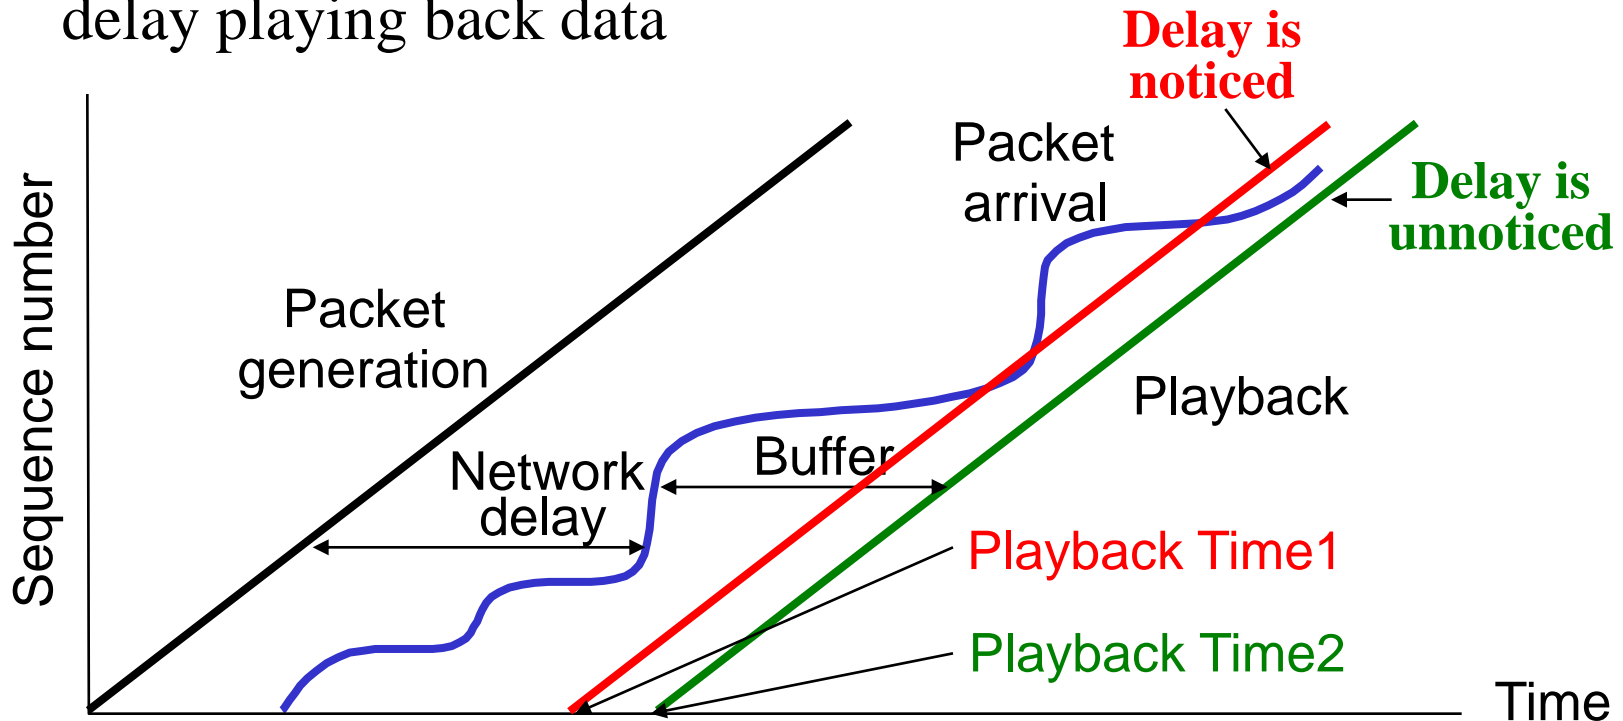
Application Requirements (Real-Time)

- The data must be **played back** at some appropriate rate
 - Each sample has a particular **playback time** in the receiver
- If data arrives **after** its appropriate playback time, either delayed or dropped, it is essentially **useless**
- It is **impossible** to make sure that all samples take exactly the same amount of time to traverse the network
 - The queue lengths **vary with time** \Rightarrow so does **delays**
- To **buffer** up some amount of data in reserve, thereby always providing a store of packets waiting to be played back




Application Requirements (Real-Time)

- As long as the playback line is far enough to the right in time, **the variation in network delay** is never noticed by the application
- For audio applications, there are limits to **how far** we can delay playing back data



Application Requirements (Real-Time)

- We can measure the **one-way delay** over a certain path across the Internet
 - The key factor is the **variability** of the delay
- Set the playback point to **100 ms** \Rightarrow 3% packets arrive too late
- Set the playback point to **200 ms** \Rightarrow ensures that all packets arrived in time



Approaches to QoS Support

- Two broad categories of approaches that support QoS:
 - **Fine-grained approaches:** provide QoS to **individual** applications or flows
 - **RSVP (Resource Reservation Protocol)**
 - **Coarse-grained approaches:** provide QoS to large classes of data or **aggregated** traffic
 - **Differentiated Services**

RSVP (Resource Reservation Protocol)

Integrated Services (RSVP)

- Service classes:
 - **Guaranteed service:** the network should guarantee that the **maximum packet delay** has some specified value
 - The application can set the **playback point** so that no packet will ever arrive after its playback time
 - **Controlled load service: tolerant, adaptive** applications
 - Uses a **queuing mechanism** to isolate the controlled load traffic from other traffic
 - Uses some form of **admission control** to limit the total amount of controlled load traffic such that the load is kept reasonable low
 - **Adjusts the playback point** as network delay varies, and controls a reasonable packet loss rate


Integrated Services (RSVP)

- Four major parts for providing integrated services:
 - **Flowspec**: the set of information that a user provides to the network regarding the service
 - **Admission control**: the process of deciding when to **say no** to a **requesting user**
 - **Resource reservation**: the mechanism by which the users of the network and the components of the network **exchange information**, such as requests for service
 - **Packet scheduling**: the mechanism of managing the way packets are **queued** and **scheduled for transmission** in the switches and routers

Flowspec

- There are two separable parts to the flowspec:
 - **RSpec**: describes the **service requested** from the network
 - For a controlled load service: **no additional parameters**
 - For a guaranteed service: specify a **delay** target or bound
 - **TSpec**: describes the flow's **traffic characteristics**
 - Specify the **requesting bandwidth**: the **average** bit rate and the **peak** bit rate

The same average bit rate
Different peak bit rates



Admission Control

- Admission control looks at the **TSpec** and **RSpec** of the flow and tries to decide if the desired service can be provided
 - Given the currently **available resources**
 - Without causing any **previously admitted** flow to receive **worse service** than it had requested
- Admission control is very dependent on the **type of requested service** and on the **queuing discipline** employed in the routers
 - For a **guaranteed service**: a good algorithm to make a **definitive yes/no decision** is required
 - For a **controlled load service**: it may be based on **heuristics**
 - E.g., current delays are **far inside the bounds** \Rightarrow It should be able to admit another flow without difficulty

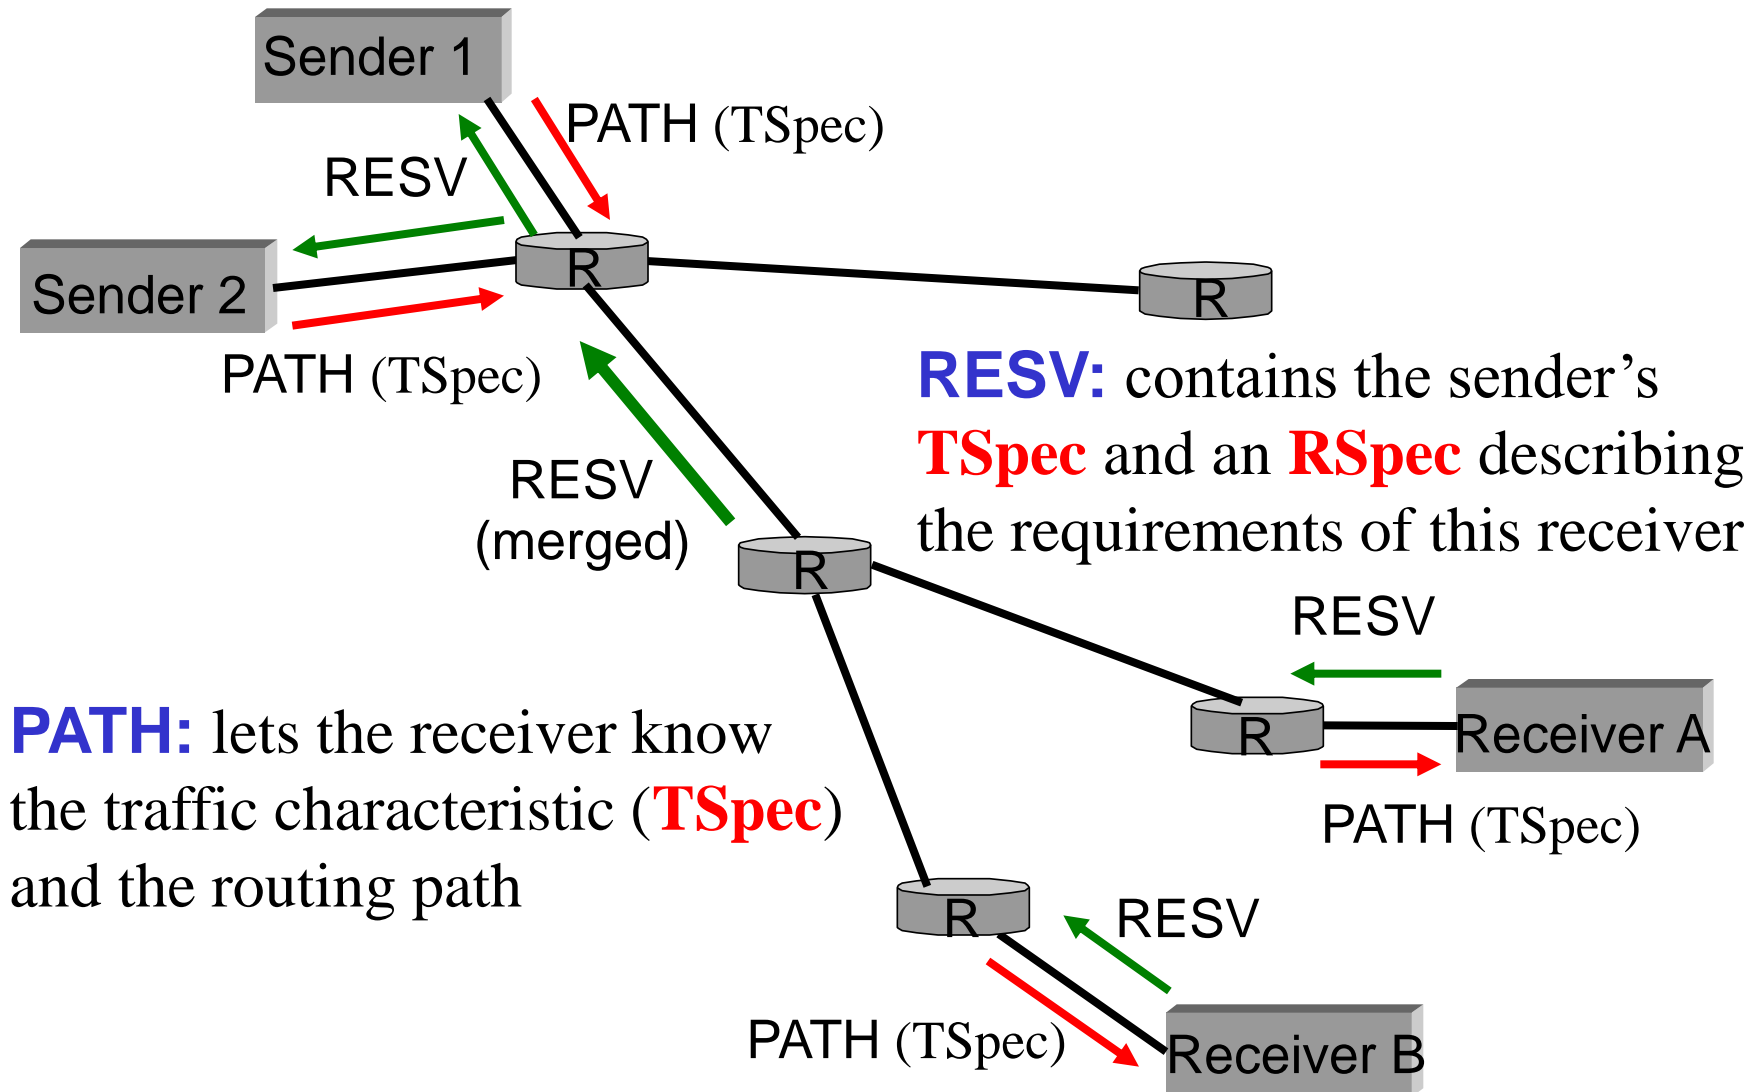
Admission Control

- **Policing:** is a function applied on a **per-packet basis** to make sure that a flow conforms to the **TSpec** that was used to make the reservation
 - There are several options, the obvious one being to **drop offending packets**
 - Another option is to **drop the offending packets first** if any packets are needed to drop

Resource Reservation (RSVP)

- The most important protocol is **Resource Reservation Protocol (RSVP)**
 - Maintains the **robustness** of connectionless networks by using the idea of **soft state** in the routers
 - Supports **multicast** flows just as effectively as **unicast** flows
- Consider the case of one sender and one receiver trying to get a reservation for the traffic flowing between them
 - The receiver needs to know **what traffic** the sender is likely to send (to make an appropriate reservation)
 - It needs to know **what path** the packets will follow (to establish a reservation **at each router on the path**)

Resource Reservation (RSVP)



Packet Scheduling

- For the routers to actually deliver the packets, all the routers on the path must
 - **Classifying packets:** **classify** each packet with the appropriate **reservation** so that it can be handled correctly
 - **Packet scheduling:** **manage** the packets in the **queues** according to the requested service
- Classifying packets is done by examining up to **five fields** in the packet: **the source address, destination address, protocol number, source port, and destination port**
 - Based on this information, the packet can be placed in the appropriate class
- Packet scheduling is an area where implementers can try to do creative things to **realize the service model efficiently**

Scalability Issues

- In the **best-effort service** model, routers store **little or no state** about the individual flows passing through them
- RSVP raises the possibility that **every flow** passing through a router might have **a corresponding reservation**
 - Each of those reservations needs some amount of state
 - Stored in memory and refreshed periodically
 - The router needs to **classify, police, and queue** each of those flows
- Suppose that every flow on an OC-48 (2.5 Gbps) link represents a 64-Kbps audio stream
 - $2.5 \times 10^9 / 64 \times 10^3 = 39,000$ flows
 - Maintaining **per-flow** state may be **not practical**

⇒ **Differentiated services**

Differentiated Services

Differentiated Services


- **Differentiated services model (DiffServ)** allocates resources to a small number of classes of traffic
 - Some proposed approaches divide traffic into **two classes**
 - Enhance the best-effort service model by adding just one new class called **“premium”**
- Each router should figure out which packets are **premium** and which are regular old **best-effort**
 - Could be done by using a bit in the packet header
 - The **router at the edge** of an Internet service provider’s network might set this bit for some packets
- In practice: based on the behavior of individual routers called **“per-hop behaviors” (PHBs)**

Differentiated Services (EF)

- **Six bits** (taken from the old **TOS** (type of service) **byte** of the IP header) have been allocated for **DiffServ code points (DSCP)**
 - Identify a particular PHB to be applied to a packet
- One of the simplest PHBs is **“expedited forwarding” (EF)**
 - The packets should be forwarded by the router with **minimal delay and loss**
 - One possible strategy is to give EF packets **strict priority** over all other packets
 - Another is to perform **weighted fair queuing**
 - The weight of EF set is sufficiently high
 - All EF packets can be delivered quickly

Differentiated Services (AF)

- Another PHB is known as **“assured forwarding” (AF)**
 - Based on an approach known as **“RED with In and Out” (RIO)** or **“Weighted RED”** (RED, random early detection)
 - A custom is allowed to send up to **y Mbps** of assured traffic
 - If it sends **less** than y Mbps, all packets are marked **“in”**
 - If it **exceeds** the rate, the excess packets are marked **“out”**



Differentiated Services (WFQ)

- A third way is **weighted fair queuing (WFQ)**: it uses the DSCP value to determine which queue to put a packet into
 - Define the **best-effort queue** and the **premium queue**
 - Choose weights that makes the premium packets get a better service than the best-effort packets
- E.g. premium queue weight = **1**, best-effort queue weight = **4**
 - The bandwidth available to premium packets is
$$B_{\text{premium}} = W_{\text{premium}} / (W_{\text{premium}} + W_{\text{best_effort}}) = 0.2$$
 - Reserve **20%** bandwidth of the link for premium packets
 - If the offered load of premium traffic is 10% of the link on average, the service of premium traffic will be very good